

# **Tutorial**

## Synthetic Smart Meter Data Generation using Variational Autoencoders

Kutay Bölat  
TU Delft

# Contents

---

- Generative modelling (reminder)
  - Latent variable models
  - Variational autoencoders
    - Hands-on session




# Generative modelling (reminder)

# Generative modelling (reminder)

$$p_{\text{real-life}}(\mathbf{x}) \rightarrow p_{\text{data}}(\mathbf{x}) = \frac{1}{N} \sum_{\mathcal{X}} \delta(\mathbf{x} - \mathbf{x}_i)$$


$$\operatorname{argmin}_{\theta} D_{KL}(p_{\text{data}}(\mathbf{x}) \| p_{\theta}(\mathbf{x})) = \operatorname{argmax}_{\theta} \frac{1}{N} \sum_{\mathcal{X}} \log p_{\theta}(\mathbf{x}_i)$$



# Generative modelling (reminder)

**Goal:** Find a (parameterised) probabilistic model  $p(\mathbf{x})$ , where  $\mathbf{x}$  is high-dimensional.

**Problem:** Finding/learning relations between many features is exceedingly hard (even for very deep and wide neural networks).



$$\begin{aligned} p(\mathbf{x}) &= p(x_1, x_2, x_3, \dots, x_d) \\ &= p(x_1)p(x_2|x_1)p(x_3|x_1, x_2) \dots p(x_d|x_1, \dots, x_{d-1}) \end{aligned}$$



# Latent variable models

# Latent variable models

**Idea:** Introduce unobserved auxiliary variables (latent variables)



- Observables:  $\mathbf{x} \in \mathbb{R}^d$
- Latent variables:  $\mathbf{z} \in \mathbb{R}^m$

**Simplified model structure:**


- $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \int p(\mathbf{z}) \prod_{j=1}^d p(x_j | \mathbf{z}) d\mathbf{z}$

**Motivation:**

- 'Meaningful' data tends to reside on a lower dimensional manifold in high-dimensional space.
- Often, known causal relations can simplify the model.




[scikit-learn.org/stable/auto\\_examples/manifold/plot\\_compare\\_methods.html](https://scikit-learn.org/stable/auto_examples/manifold/plot_compare_methods.html)




<https://www.frontiersin.org/articles/10.3389/fpsyg.2017.00798/full>

# Latent variable models

- Gaussian mixture models



Requires  $p(\mathbf{x})$  to be expressive enough to capture non-linearities.



$$p(\mathbf{x}) = \sum_z p(\mathbf{x}|z)p(z) = \sum_{z \in \{0,1,2\}} \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_z)p(z)$$







# Variational autoencoders

# Variational autoencoders - Introduction

- Introduced in Kingma, D. P., & Welling, M. (2013). Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*
  - Over 37000 citations!
- Two interpretations:
  - A probabilistic autoencoder
  - Latent variable models parametrized by NNs



# Variational autoencoders – End-result



## Intuitive objective:

1. Good reconstruction:  $\tilde{x}_i \approx x_i$  ( $z_i$  holds important information about  $x_i$ )
2. Regularization:  $p(z_i|x_i) \approx N(0, I)$  ( $z_i$  holds important information about other data points)

$$\mathcal{L}(\phi, \theta, X) = \underbrace{-\frac{1}{NM} \sum_{i,j} \left\| \mathbf{x}^{(i)} - \mu_\theta(g_\phi(\mathbf{x}^{(i)}, \epsilon_j^{(i)})) \right\|_2^2}_{\text{reconstruction error}} \underbrace{- \frac{1}{2} \sum_{k=1}^m \left( \mu_{k\phi}(\mathbf{x}^{(i)}) \right)^2 + \left( \sigma_{k\phi}(\mathbf{x}^{(i)}) \right)^2 - 1 - 2 \log \sigma_{k\phi}(\mathbf{x}^{(i)})}_{\text{regularization}}$$

# Latent variable models - Math

## Definitions

- The model is  $p_\theta(\mathbf{x})$  where  $\int_{\mathbf{x}} p_\theta(\mathbf{x}) d\mathbf{x} = 1$  and  $p_\theta(\mathbf{x}) > 0$ .
- $p_\theta(\mathbf{x}) = \int_{\mathbf{z}} p_\theta(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \int_{\mathbf{z}} p_\theta(\mathbf{x}|\mathbf{z}) p_\theta(\mathbf{z}) d\mathbf{z}$ .

## Some useful operations

- $p_\theta(\mathbf{x}|\mathbf{z})$ : Called the **likelihood** or the **inference** model. Very flexible. Can be modelled with a neural network.
- $p_\theta(\mathbf{z})$ : Called the **prior**. Can be modelled with simple distributions and it transfers prior assumptions to the model.
- $\int_{\mathbf{z}} \dots d\mathbf{z}$ : Averaging over latent variables. **Marginalization**. Generally **intractable**, i.e. impossible to compute analytically and extremely expensive to compute numerically.
- Also, we might want to infer (low dimensional) latent variables by looking at (high dimensional) data, i.e. **compression**. This requires us to compute the **posterior**

$$p_\theta(\mathbf{z}|\mathbf{x}) = \frac{p_\theta(\mathbf{x}, \mathbf{z})}{p_\theta(\mathbf{x})} = \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\theta(\mathbf{z})}{\int_{\mathbf{z}} p_\theta(\mathbf{x}, \mathbf{z}) d\mathbf{z}}$$

which is also intractable most of the time.

- With  $p_\theta(\mathbf{z})$  and  $p_\theta(\mathbf{x}|\mathbf{z})$  we can generate data!
- ... but learning  $p_\theta(\mathbf{x}|\mathbf{z})$  requires  $p_\theta(\mathbf{z}|\mathbf{x})$
- We will use an **approximate posterior**  $q_\phi(\mathbf{z}|\mathbf{x})$  (a separate model)

# Variational autoencoders – Objective

- Maximum log-likelihood criteria:  $\operatorname{argmax}_{\theta} \mathbb{E}_{p_{data}(\mathbf{x})} \log p_{\theta}(\mathbf{x}) = \operatorname{argmax}_{\theta} \frac{1}{N} \sum_{i=1}^N \log p_{\theta}(\mathbf{x}_i)$
- $\log p_{\theta}(\mathbf{x}) = \log \left( \frac{p_{\theta}(\mathbf{x}|\mathbf{z})p_{\theta}(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right) = \log \left( \frac{p_{\theta}(\mathbf{x}|\mathbf{z})p_{\theta}(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right) = \log(p_{\theta}(\mathbf{x}|\mathbf{z})) - \log \left( \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z})} \right) + \log \left( \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right)$
- $\log p_{\theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z})) - \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log \left( \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z})} \right) + \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log \left( \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right)$   
 $= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z})) - D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}) \right) + D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}|\mathbf{x}) \right)$ 
  - $\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z}))$ : Reconstruction log-likelihood
  - $D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}) \right)$ : KL-divergence (discrepancy) between the approx. posterior and the prior
  - $D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}|\mathbf{x}) \right)$ : KL-divergence (discrepancy) between the approx. posterior and the true posterior

# Variational autoencoders – Objective

$$\log p_{\theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z})) - D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}) \right) + D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}|\mathbf{x}) \right)$$

- Remarks:
  - KL-divergence is always positively defined.
  - The only intractable term is  $D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}|\mathbf{x}) \right)$ .
- $\log p_{\theta}(\mathbf{x}) \geq \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z})) - D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}) \right) = ELBO$
- Objective:

$$\phi^*, \theta^* = \operatorname{argmax}_{\phi, \theta} \mathbb{E}_{p_{data}(\mathbf{x})} \left[ \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z}))}_{\text{reconstruction success}} - \underbrace{D_{KL} \left( q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}) \right)}_{\text{KL-divergence}} \right]$$

# Variational autoencoders - Recipe

- How to implement and train VAEs?
- **Step 1:** Selection of distributions
  - $p_{\theta}(\mathbf{x}|\mathbf{z})$ : Totally depends on the application. It can be:
    - $p_{\theta}(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_{\theta}(\mathbf{z}), \boldsymbol{\Sigma}_{\theta}(\mathbf{z}))$  if  $\mathcal{X} = \mathbb{R}^d$
    - $p_{\theta}(\mathbf{x}|\mathbf{z}) = \text{Lognormal}(\mathbf{x}|\boldsymbol{\mu}_{\theta}(\mathbf{z}), \boldsymbol{\Sigma}_{\theta}(\mathbf{z}))$  if  $\mathcal{X} = \mathbb{R}^{+d}$
    - $p_{\theta}(\mathbf{x}|\mathbf{z}) = \text{Beta}(\mathbf{x}|\boldsymbol{\beta}_{\theta}(\mathbf{z}), \boldsymbol{\alpha}_{\theta}(\mathbf{z}))$  if  $\mathcal{X} = [0,1]^d$
    - $p_{\theta}(\mathbf{x}|\mathbf{z}) = \text{Categorical}(\mathbf{x}|\mathbf{P}_{\theta}(\mathbf{z}))$  if  $\mathcal{X} = \{0,1, \dots, K\}^d$
  - $q_{\phi}(\mathbf{z}|\mathbf{x})$ : Should result in an analytical  $D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}))$  expression. Also, we must be able to take samples in a “straightforward” manner. (More on this later)
  - $p_{\theta}(\mathbf{z})$ : Should result in an analytical  $D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z}))$  expression.

# Variational autoencoders - Recipe

- Investigating the ELBO – **Reconstruction log-likelihood**

- $$\mathbb{E}_{p_{data}(\mathbf{x})} \left[ \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log(p_{\theta}(\mathbf{x}|\mathbf{z})) \right] = \frac{1}{N} \sum_{i=1}^N \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \log(p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z})) \quad (\text{No closed form expression})$$

$$= \frac{1}{NM} \sum_{i=1}^N \sum_{j=1}^M \log(p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}_j^{(i)})), \quad \text{where } \mathbf{z}_j^{(i)} \sim q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}) \quad (\text{Monte Carlo approximation})$$

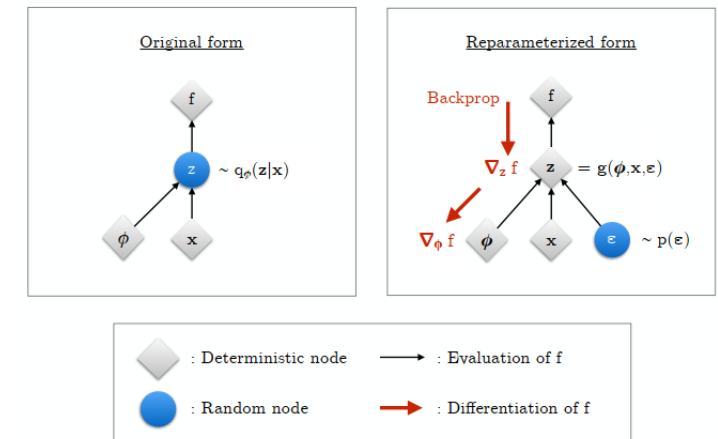
- Problem:** Sampling is not a differentiable operation.

- Solution:** Reparameterization trick (**Step 2**)

- $\mathbf{z}_j^{(i)} \sim q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}) \rightarrow \mathbf{z}_j^{(i)} = g_{\phi}(\mathbf{x}^{(i)}, \epsilon_j^{(i)}), \quad \text{where } \epsilon_j^{(i)} \sim p(\epsilon)$
- This separation should be “straightforward”, i.e. computationally cheap.

- For  $q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{\phi}(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}_{\phi}(\mathbf{x})))$

- $\mathbf{z}_j^{(i)} = \boldsymbol{\mu}_{\phi}(\mathbf{x}^{(i)}) + \boldsymbol{\sigma}_{\phi}(\mathbf{x}^{(i)}) \odot \boldsymbol{\epsilon}_j^{(i)}, \quad \text{where } \boldsymbol{\epsilon}_j^{(i)} \sim \mathcal{N}(\boldsymbol{\epsilon}|\mathbf{0}, \mathbf{I})$



Kingma, Diederik P., and Max Welling. "An Introduction to Variational Autoencoders." *Foundations and Trends® in Machine Learning*, vol. 12, no. 4, 2019, pp. 307–92. Crossref, <https://doi.org/10.1561/22000000056>.




# Variational autoencoders - Recipe

- Investigating the ELBO – **KL divergence**
  - Step 3:** Analytically derive the KL divergence term
  - Let's keep  $q_\phi(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_\phi(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}_\phi(\mathbf{x})))$  and  $p_\theta(\mathbf{z}) = \mathcal{N}(\mathbf{z}|\mathbf{0}, \mathbf{I})$
  - $D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\theta(\mathbf{z})) = \frac{1}{2} \sum_{k=1}^m \left( \mu_{k\phi}(\mathbf{x}) \right)^2 + \left( \sigma_{k\phi}(\mathbf{x}) \right)^2 - 1 - 2 \log \sigma_{k\phi}(\mathbf{x})$
- Step 3.5:** Put everything together
- For  $p_\theta(\mathbf{x}^{(i)} | g_\phi(\mathbf{x}^{(i)}, \boldsymbol{\epsilon}_j^{(i)})) = \mathcal{N}(\mathbf{x}^{(i)} | \boldsymbol{\mu}_\theta(g_\phi(\mathbf{x}^{(i)}, \boldsymbol{\epsilon}_j^{(i)})), \mathbf{I})$ , ELBO becomes
  - $$\mathcal{L}(\phi, \theta, X) = \underbrace{-\frac{1}{NM} \sum_{i,j} \left\| \mathbf{x}^{(i)} - \boldsymbol{\mu}_\theta(g_\phi(\mathbf{x}^{(i)}, \boldsymbol{\epsilon}_j^{(i)})) \right\|_2^2}_{\text{reconstruction error}} \underbrace{- \frac{1}{2} \sum_{k=1}^m \left( \mu_{k\phi}(\mathbf{x}^{(i)}) \right)^2 + \left( \sigma_{k\phi}(\mathbf{x}^{(i)}) \right)^2 - 1 - 2 \log \sigma_{k\phi}(\mathbf{x}^{(i)})}_{\text{KL-divergence}}$$

# Variational autoencoders - Recipe

- **Step 4:** Replace all the parameterized functions with neural networks



$$\mathcal{L}(\phi, \theta, X) = -\frac{1}{NM} \sum_{i,j} \left\| \mathbf{x}^{(i)} - \boldsymbol{\mu}_\theta(g_\phi(\mathbf{x}^{(i)}, \boldsymbol{\epsilon}_j^{(i)})) \right\|_2^2 - \frac{1}{2} \sum_{k=1}^m \left( \mu_{k\phi}(\mathbf{x}^{(i)}) \right)^2 + \left( \sigma_{k\phi}(\mathbf{x}^{(i)}) \right)^2 - 1 - 2 \log \sigma_{k\phi}(\mathbf{x}^{(i)})$$

- **Step 5:** Let your favourite DL library (PyTorch, Tensorflow) does its magic.

# Variational autoencoders - Recipe

- Investigating the ELBO – **Overall**


- $$\mathcal{L}(\phi, \theta, X) = -\frac{1}{NM} \sum_{i,j}^{n,M} \left\| \mathbf{x}^{(i)} - \boldsymbol{\mu}_{\theta} \left( g_{\phi} \left( \mathbf{x}^{(i)}, \boldsymbol{\epsilon}_j^{(i)} \right) \right) \right\|_2^2 -$$

$$\frac{1}{2} \sum_{k=1}^m \left( \mu_{k\phi}(\mathbf{x}^{(i)}) \right)^2 + \left( \sigma_{k\phi}(\mathbf{x}^{(i)}) \right)^2 - 1 - 2 \log \sigma_{k\phi}(\mathbf{x}^{(i)})$$
- $-\mathcal{L}(\phi, \theta, X)$  is pretty much identical to "Autoencoder/Reconstruction Loss" + "KL regularization"
- Maximizing ELBO (minimizing  $-\mathcal{L}(\phi, \theta, X)$ ) encourages
  - Inputs and outputs to be the same as much as possible
  - $\mu_{k\phi}(\mathbf{x}^{(i)})$ s to be close to 0 as much as possible
  - $\sigma_{k\phi}(\mathbf{x}^{(i)})$ s to be close to 1 as much as possible

# Variational autoencoders - Recipe

- **An important remark**


- We use  $p(z)$  for two reasons in practice:
  1. Regularizing the latent encodings **during training**.
  2. Generating new data **during inference**.




- For generating synthetic data, we DO NOT have any access to original data or their latent encodings!

# Variational autoencoders - Examples

- Image data (MNIST)



Original




Reconstruction

[github.com/clementchadebec/benchmark\\_VAE](https://github.com/clementchadebec/benchmark_VAE)

# Variational autoencoders - Examples

- Image data (MNIST)



Samples

[github.com/clementchadebec/benchmark\\_VAE](https://github.com/clementchadebec/benchmark_VAE)

# Variational autoencoders - Examples

- Image data (CelebA)



Original




Reconstruction

[github.com/clementchadebec/benchmark\\_VAE](https://github.com/clementchadebec/benchmark_VAE)

# Variational autoencoders - Examples

- Image data (CelebA)



Samples



[github.com/clementchadebec/benchmark\\_VAE](https://github.com/clementchadebec/benchmark_VAE)



# Variational autoencoders - Extensions

- Conditional VAE

- $p_{\theta}(x) \rightarrow p_{\theta}(x|y)$
- ELBO:  $\mathbb{E}_{q_{\phi}(z|x, y)} \log(p_{\theta}(x|z, y)) - D_{KL}(q_{\phi}(z|x, y) \parallel p_{\theta}(z, y))$





[ijdykeman.github.io/ml/2016/12/21/cvae.html](http://ijdykeman.github.io/ml/2016/12/21/cvae.html)



# Hands-on session

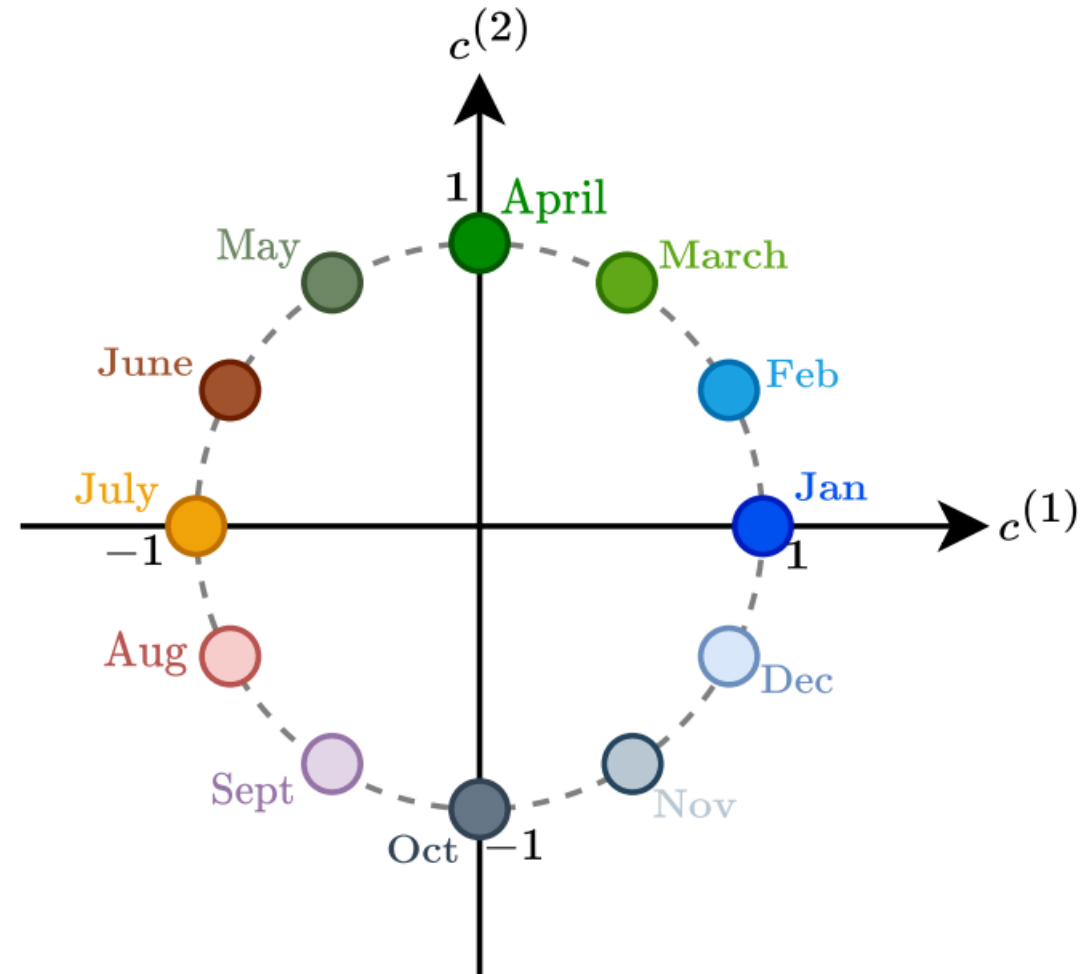
# Hands-on session

- Data representation
  - Daily snapshots




# Hands-on session

- Conditions
  - Months and days of the week
  - Circular transform




# Hands-on session



# Hands-on session



- [github.com/kabolat/innocypes-summer-school\\_synthetic-data-tutorial](https://github.com/kabolat/innocypes-summer-school_synthetic-data-tutorial)



# Thanks for your attention!

## Any questions?

Kutay Bölat

 K.Bolat@tudelft.nl  
 [github.com/kabolat](https://github.com/kabolat)



This project has received funding from the European Union's Horizon 2020 research and innovation programme under the Marie Skłodowska-Curie grant agreement No 956433.

